

# A brief introduction to OpenMP

## 2: Data sharing clauses



- ▶ Recap
- ▶ Data sharing clauses
- ▶ The Pi program
- ▶ Critical regions
- ▶ False sharing issues
- ▶ Reductions

- ▶ Fork/join execution model.
- ▶ Shared memory model:
  - ▶ All threads can read/write the *same* memory.
- ▶ Set number of threads with OMP\_NUM\_THREADS environment variable.
- ▶ Parallelise simple loops with worksharing clauses:

```
#pragma omp parallel for
for (int i = 0; i < N; ++i) {
    A[i] = ...;
}
```

- ▶ Talked about collapse, nowait and schedule clauses.

Remember: OpenMP is a *shared memory* programming model.

- ▶ By default, all data is available to all threads.
- ▶ There is a single copy of *shared* data.

You must specify which data should be *private* to each thread.

- ▶ Each thread then has local (stack) space for each private variable.
- ▶ Each copy is only visible to its associated thread.

## Notice

Variables declared *inside* the parallel region will be private to each thread.

- ▶ All data on the heap is shared.
- ▶ Therefore all the data allocated with `malloc` is shared.
- ▶ You must ensure that different threads do not write to the same element of these arrays.

## Caution

Setting a data sharing clause on a heap variable only effects the metadata of the variable. The pointer could be private, but the target will still be shared.

- ▶ `shared(x)` There is one copy of the `x` variable. The programmer must ensure synchronisation.
- ▶ `private(x)` Each thread gets its own local `x` variable. It is not initialised. The value of the original `x` variable is undefined on region exit.
- ▶ `firstprivate(x)` Each thread gets its own `x` variable, and it is initialised to the value of the original variable entering the region.
- ▶ `lastprivate(x)` Used for loops. Each thread gets its own `x` variable, and on exiting the region the original variable is updated taking the value from the sequentially last iteration.

These are the most common clauses that are needed.

Simple **for** loop, which just sets a variable to the iteration number. Each iteration prints out the current and next value of *x*, along with the thread number. Will see what happens with different data sharing clauses.

```
1  int x = -1;
2  #pragma omp parallel for private(x) / firstprivate(x) /
   ↪ lastprivate(x)
3  for (int i = 0; i < N; ++i) {
4      printf("Thread %d setting x=%d to %d\n",
   ↪ omp_get_thread_num(), x, i);
5      x = i;
6  }
```

*N* is set to 10. Ran using 4 threads.

```
private:
  before: x=-1
    Thread 1 setting x=0 to 3
    Thread 2 setting x=0 to 6
    Thread 3 setting x=0 to 8
    Thread 0 setting x=0 to 0
    Thread 1 setting x=3 to 4
    Thread 2 setting x=6 to 7
    Thread 3 setting x=8 to 9
    Thread 0 setting x=0 to 1
    Thread 1 setting x=4 to 5
    Thread 0 setting x=1 to 2
  after: x=-1
```

Each thread starts with its own x. No guarantees of initial value, but happened to be zero this time.



```
firstprivate:
  before: x=-1
    Thread 3 setting x=-1 to 8
    Thread 2 setting x=-1 to 6
    Thread 1 setting x=-1 to 3
    Thread 0 setting x=-1 to 0
    Thread 3 setting x=8 to 9
    Thread 2 setting x=6 to 7
    Thread 1 setting x=3 to 4
    Thread 0 setting x=0 to 1
    Thread 1 setting x=4 to 5
    Thread 0 setting x=1 to 2
  after: x=-1
```

Each thread starts with its own x, which set to the value of x before entering the parallel region, -1.

```
lastprivate:  
  before: x=-1  
    Thread 3 setting x=2 to 8  
    Thread 2 setting x=1 to 6  
    Thread 1 setting x=0 to 3  
    Thread 3 setting x=8 to 9  
    Thread 0 setting x=-1 to 0  
    Thread 2 setting x=6 to 7  
    Thread 1 setting x=3 to 4  
    Thread 0 setting x=0 to 1  
    Thread 1 setting x=4 to 5  
    Thread 0 setting x=1 to 2  
  after: x=9
```


Each thread starts with its own  $x$ , which set to to a garbage value.  
On exiting the region, the original  $x$  is set to the value of the last iteration of the loop, 9.

Use a simple program to numerically approximate  $\pi$  to explore:

- ▶ Use of data sharing clauses.
- ▶ Updating a shared variable in parallel.
- ▶ Reductions.


# Integration to calculate Pi

$$\int_0^1 \frac{4}{1+x^2} dx = \pi$$



# Trapezoidal rule

Sum the area of the boxes. Choose a small *step* size to generate lots of boxes, and increase accuracy.



We will use this code which calculates the value of  $\pi$  as an example for the remainder of this session.

```
1  double step, x, sum, pi;
2  step = 1.0/num_steps;
3  for (int ii = 1; ii <= num_steps; ++ii) {
4      x = (ii-0.5)*step;
5      sum = sum + (4.0/(1.0+x*x));
6  }
7  pi = step * sum;
```

With 100,000,000 steps, this takes 0.368s on my laptop.

Use a worksharing directive to parallelise the loop.

```
1  double step, x, sum, pi;  
2  step = 1.0/num_steps;  
3  #pragma omp parallel for private(x)  
4  for (int ii = 1; ii <= num_steps; ++ii) {  
5      x = (ii-0.5)*step;  
6      sum = sum + (4.0/(1.0+x*x));  
7  }  
8  pi = step * sum;
```

What about data sharing?

- ▶ `x` needs to be used independently by each thread, so mark as `private`.
- ▶ `sum` needs to be updated by *all* threads, so leave as `shared`.

- ▶ But need to be careful changing the shared variable, sum.
- ▶ All threads can update this value directly!
- ▶ A critical region only allows one thread to execute at any one time. No guarantees of ordering.

```
1  double step, x, sum, pi;
2  step = 1.0/num_steps;
3  #pragma omp parallel for private(x)
4  for (int ii = 1; ii <= num_steps; ++ii) {
5      x = (ii-0.5)*step;
6      #pragma omp critical
7      {
8          sum = sum + (4.0/(1.0+x*x));
9      }
10 }
11 pi = step * sum;
```



Run on a MacBook Pro (Intel Core i7-4980HQ CPU @ 2.80GHz)  
with 4 threads.

Implementation	Runtime (s)
Serial	0.368
Critical	426.1
Atomic	8.3

Slower than serial!

- ▶ Criticals or atomics methods cause threads to synchronise for every update to `sum`.
- ▶ But each thread could compute a partial sum independently, synchronising once to total at the end.

Make `sum` an array of length equal to the number of threads.

- ▶ Each thread stores its partial sum, and the array is totalled by the master thread serially at the end.
- ▶ As it's *shared memory*, the `sum` array can be read just fine on the master rank.

```
1  step = 1.0/num_steps;
2  #pragma omp parallel private(x,tid)
3  {
4  tid = omp_get_thread_num();
5  sum[tid] = 0.0;
6  #pragma omp for
7  for (int ii = 1; ii <= num_steps; ++ii) {
8      x = (ii-0.5)*step;
9      sum[tid] = sum[tid] + (4.0/(1.0+x*x));
10 }
11 }
12 for (int ii = 0; ii < nthreads; ++ii) {
13     pi = pi + sum[ii];
14 }
15 pi = pi * step;
```

Run on a MacBook Pro (Intel Core i7-4980HQ CPU @ 2.80GHz)  
with 4 threads.

Implementation	Runtime (s)
Serial	0.368
Critical	426.1
Atomic	8.3
Array	2.8*

\* Might be faster if have a smart compiler which avoided false sharing.


Fastest parallel version so far, but still slow.

# False sharing



This code is susceptible to *false sharing*.

- ▶ False sharing occurs when different threads update data on the same cache line.
- ▶ It is a *different* phenomenon to cache thrashing, as result of parallel shared memory execution.
- ▶ Cache system is coherent between cores, so data consistency must be maintained.
- ▶ The cache line is no longer up to date because another thread changed it (in their local cache).
- ▶ Therefore, cache line must be flushed to memory and reread into the other thread every time.



Can use data sharing clauses to our advantage here:

Give each thread a *scalar* copy of *sum* to compute their partial sum, and reduce with only one critical (or atomic) region at the end. No false sharing, as value is just a single number (i.e. a register).

```
1  step = 1.0/num_steps;
2  #pragma omp parallel private(x) firstprivate(sum)
3  {
4  #pragma omp for
5  for (int ii = 1; ii <= num_steps; ++i) {
6      x = (ii-0.5)*step;
7      sum = sum + (4.0/(1.0+x*x));
8  }
9  #pragma omp critical
10 pi = pi + sum;
11 }
12 pi = pi * step;
```

Run on a MacBook Pro (Intel Core i7-4980HQ CPU @ 2.80GHz)  
with 4 threads.

Implementation	Runtime (s)
Serial	0.368
Critical	426.1
Atomic	8.3
Array	2.8
First private	0.104

Finally faster than serial! Around 3.5X faster on 4 threads.

Much simpler to use the OpenMP reduction clause on a worksharing loop. Specify the operation and the variable.

- |  |                                     |
|--|-------------------------------------|
| ▶ <code>reduction(+:var)</code>          | ▶ <code>reduction(^:var)</code>     |
| ▶ <code>reduction(-:var)</code>          | ▶ <code>reduction(&amp;:var)</code> |
| ▶ <code>reduction(*:var)</code>          | ▶ <code>reduction( :var)</code>     |
| ▶ <code>reduction(&amp;&amp;:var)</code> | ▶ <code>reduction(min:var)</code>   |
| ▶ <code>reduction(  :var)</code>         | ▶ <code>reduction(max:var)</code>   |

Can also do array reductions. Each element of array is treated as own, separate, reduction. Similar to:

```
MPI_Allreduce(MPI_IN_PLACE, arr, N, MPI_DOUBLE,  
  ↪ MPI_SUM, 0, MPI_COMM_WORLD);
```



Much simpler to write using the reduction clause — just need a single directive:

```
1  step = 1.0/num_steps;
2  #pragma omp parallel for private(x)
   ↪ reduction(+:sum)
3  for (int ii = 1; ii <= num_steps; ++i) {
4      x = (ii-0.5)*step;
5      sum = sum + (4.0/(1.0+x*x));
6  }
7  pi = step * sum;
```

Run on a MacBook Pro (Intel Core i7-4980HQ CPU @ 2.80GHz)  
with 4 threads.

Implementation	Runtime (s)
Serial	0.368
Critical	426.1
Atomic	8.3
Array	2.8
First private	0.104
Reduction	0.095

Around 3.9X faster on 4 threads!

## Recommendation

Use the reduction clause for reductions.

- ▶ Have now covered the most common parts of OpenMP.
- ▶ 80/20 rule: Most programs will only use what you know so far.
- ▶ OpenMP is deceptively simple!